

Mutual Exclusion

CS511

Example I: Undesired Interleavings

Consider two processes that increment a counter and print the result

```
1 global int x=0;
2
3 thread P: {
4     x = x + 1;
5     print(x);
6 }
7
8 thread Q: {
9     x = x + 1;
10    print(x);
11 }
```

What output do you expect? Let's run it and see what happens...

Atomicity Assumption on Assignments

- ▶ An instruction such as $x=x+1$ is actually compiled into simpler sets of assembly instructions

```
1 load x,r
2 add r,1
3 write r,x
```

- ▶ Could cause undesired interleavings at assembly level

```
1 global int x=0;
2
3 thread P: {
4     x = x + 1;
5     print(x);
6 }
7
8 thread Q: {
9     x = x + 1;
10    print(x);
11 }
```

Race Condition

Race Condition

A timing dependent error involving shared state

- ▶ Whether it happens depends on how threads are scheduled
- ▶ Once thread T1 starts doing something, it needs to “race” to finish it because if thread T2 looks at the shared variable before T1 is done, it may see something inconsistent
- ▶ Hard to detect

Atomicity Assumption on Assignments

Atomic Operation

An operation is atomic if it is executed until completion without interleaving statements from other processes

- ▶ Atomic operations are the smallest unit in which a trace can be listed
- ▶ We assume throughout this course that assignment statements are **atomic** for scalar values (i.e. numbers, booleans, chars, etc.)
- ▶ Eg. $x=2$ is atomic but $x=x+1$ is not.

Atomicity Assumption on Assignments

- ▶ Unfortunately, even if assignments such as $x=x+1$ were considered atomic, there are still situations where unwanted interleavings can appear
- ▶ Indeed, an easy example is obtained by “simulating” the problem with the counters using a temporary variable

Example II: Turnstile

- ▶ Consider two turnstiles to control access to a room
 - ▶ People can enter through any of them
 - ▶ People arrive at random times
 - ▶ We will assume that the room has capacity for 100 assistants
- ▶ Write a program in which the entering cycle, for each turnstile, runs in its own thread
- ▶ There should be a variable that counts the total number of people that entered

Example II: Turnstile

- ▶ Lets look at the code for one turnstile
- ▶ We assume a global variable `counter`
- ▶ Also, we assume that each turnstile only lets at most 50 people pass

```
1 global int counter=0;
2
3 thread P: {
4     for(int j = 0; j < 50; j++){
5         print("1: In comes "+ j);
6         int temp = counter+1;
7         counter = temp;
8     }
9     print("1: total " + counter);
10 }
```


Example II in Hydra

```
1 global int counter=0;
2
3 thread P: {
4     for(int j = 0; j < 50; j++){
5         print("1: In comes " + j);
6         int temp = counter+1;
7         counter = temp;
8     }
9     print("1: total " + counter);
10
11 thread Q: {
12     for(int j = 0; j < 50; j++){
13         print("2: In comes " + j);
14         int temp = counter+1;
15         counter = temp;
16     }
17     print("2: total " + counter);
18 }
```

Example II: Bad Trace

(IP Thread1, IPThread2, temp Thread1, temp Thread2, counter)

(P6, Q6, 0, 0, 0)
→ (P7, Q6, 1, 0, 0)
→ (P7, Q7, 1, 1, 0)
→ (P8, Q7, 1, 1, 1)
→ (P8, Q8, 1, 1, 1)

```
1 global int counter=0;
2
3 thread P: {
4     for(int j = 0; j < 50; j++){
5         print("1: In comes " + j);
6         int temp = counter+1;
7         counter = temp;
8     }
9     print("1: total " + counter);
10 }
```

Example II in Java

```
1 public class Example2 implements Runnable{
2     public static int counter = 0;
3     int id;
4     public Example2(int esp){ id = esp; }
5     public void run(){
6         Random rnd= new Random();
7         for(int i = 0; i < 50; i++){
8             int temp = ++counter;
9             counter = temp;
10            System.out.println(id+"- In comes: "+i);
11            try {
12                Thread.sleep(rnd.nextInt(10));
13            } catch (InterruptedException e) {
14                e.printStackTrace();
15            }
16        }
17        System.out.println(id+"- Total incoming: "+ counter);
18    }
19
20    public static void main(String args[])
21    {
22        new Thread(new Example2(1)).start();
23        new Thread(new Example2(2)).start();
24    }
25 }
```

Properties of Traces

A property is an assertion that is true for every possible trace

- ▶ Safety: A trace never reaches a “bad” state
- ▶ Liveness: Eventually, every trace will reach a “good” state

Synchronization

Mechanism that restricts the possible traces of a concurrent program in order to ensure safety and liveness properties.

- ▶ An example of a safety property follows

Race Conditions, Atomicity

Critical Sections

An Advanced Algorithm

Notions

Critical Section

A part of the program that accesses shared memory and which we wish to execute atomically

Mutual Exclusion

The problem of ensuring that two **threads** do not execute a critical section simultaneously.

```
1 thread P: {                               1 thread Q: {
2   // non-critical section                 2   // non-critical section
3   entry to critical section;              3   entry to critical section
4   // CRITICAL SECTION                    4   // CRITICAL SECTION
5   exit from critical section              5   exit from critical section
6   // non-critical section                 6   // non-critical section
7 }                                         7 }
```

Notions

Critical Section

A part of the program that accesses shared memory and which we wish to execute atomically

Mutual Exclusion

The problem of ensuring that two **threads** do not execute a critical section simultaneously.

```
1 thread P: {
2   // non-critical section
3   entry to critical section;
4   // CRITICAL SECTION
5   exit from critical section
6   // non-critical section
7 }

1 thread Q: {
2   // non-critical section
3   entry to critical section
4   // CRITICAL SECTION
5   exit from critical section
6   // non-critical section
7 }
```

Notions

Critical Section

A part of the program that accesses shared memory and which we wish to execute atomically

Mutual Exclusion

The problem of ensuring that two **threads** do not execute a critical section simultaneously.

```
1 thread P: {                               1 thread Q: {
2   // non-critical section                 2   // non-critical section
3   entry to critical section;              3   entry to critical section
4   // CRITICAL SECTION                    4   // CRITICAL SECTION
5   exit from critical section              5   exit from critical section
6   // non-critical section                 6   // non-critical section
7 }                                         7 }
```


Notions

```
1 thread P: {                               1 thread Q: {
2   // non-critical section                 2   // non-critical section
3   entry to critical section;              3   entry to critical section
4   // CRITICAL SECTION                    4   // CRITICAL SECTION
5   exit from critical section              5   exit from critical section
6   // non-critical section                 6   // non-critical section
7 }                                         7 }
```

Assumptions:

- ▶ There are no shared variables between the critical section and the non critical section (nor with the entry/exit protocol).
- ▶ The critical section always terminates.
- ▶ The scheduler is weakly fair.
 - ▶ A process that is waiting to run, will be able to do so eventually

What is Required to Guarantee Mutual Exclusion?

- ▶ In theory, mechanisms involving shared variables suffice
 - ▶ We are going to see numerous examples
 - ▶ However: Too low-level, easy to get wrong
- ▶ In practice:
 - ▶ PLs provide high-level abstractions that serve for this, and other, purposes (eg. semaphores, monitors, etc.)
 - ▶ Communication without shared variables (eg. message passing)
 - ▶ Hardware support (special atomic instructions)

The Mutual Exclusion Problem (MEP)

1. **Mutex:** At any point in time, there is at most one thread in the critical region
2. **Absence of deadlock:** If various threads try to enter the critical section, at least one of them will succeed
3. **Free from starvation:** A thread trying to enter its critical section will eventually be able to do so

General Scheme to Address the MEP

```
1 global variables;
2
3 thread {
4     non-critical section // shared variables not used
5     entry to critical section
6     CRITICAL SECTION      // may use shared variables
7     exit from critical section
8     non-critical section // shared variables not used
9 }
```

Question

Can we solve the MEP for two processes assuming that the only atomic operations are

- ▶ **read** and
- ▶ **write** on variables?

Attempt 0

```
global boolean flag = false;

thread P: {
    // non-critical section
    while (flag) {};
    flag = true;
    // CRITICAL SECTION
    flag = false;
    // non-critical section
}

thread Q: {
    // non-critical section
    while (flag) {};
    flag = true;
    // CRITICAL SECTION
    flag = false;
    // non-critical section
}
```

Attempt 0

```
global boolean flag = false;

thread P: {
    // non-critical section
    while (flag) {};
    flag = true;
    // CRITICAL SECTION
    flag = false;
    // non-critical section
}

thread Q: {
    // non-critical section
    while (flag) {};
    flag = true;
    // CRITICAL SECTION
    flag = false;
    // non-critical section
}
```

- ▶ flag flags whether someone is in the CR
- ▶ We often abbreviate `while (cond) { };` as `await (!cond)` (this is called a **busy-wait loop**)
- ▶ Let's apply this to the above example

Attempt 0

```
global boolean flag = false;

thread P: { //
    // non-critical section
    await !flag;
    flag = true;
    // CRITICAL SECTION
    flag = false;
    // non-critical section
}

thread Q: {
    // non-critical section
    await !flag;
    flag = true;
    // CRITICAL SECTION
    flag = false;
    // non-critical section
}
```

- ▶ Mutex:
- ▶ Absence deadlock:
- ▶ Free from starvation:

Attempt 0

```
global boolean flag = false;

thread P: { //
    // non-critical section
    await !flag;
    flag = true;
    // CRITICAL SECTION
    flag = false;
    // non-critical section
}

thread Q: {
    // non-critical section
    await !flag;
    flag = true;
    // CRITICAL SECTION
    flag = false;
    // non-critical section
}
```

- ▶ Mutex: No
- ▶ Absence deadlock:
- ▶ Free from starvation:

Attempt 0

```
global boolean flag = false;

thread P: { //
    // non-critical section
    await !flag;
    flag = true;
    // CRITICAL SECTION
    flag = false;
    // non-critical section
}

thread Q: {
    // non-critical section
    await !flag;
    flag = true;
    // CRITICAL SECTION
    flag = false;
    // non-critical section
}
```

- ▶ Mutex: No
- ▶ Absence deadlock: Yes
- ▶ Free from starvation:

Attempt 0

```
global boolean flag = false;

thread P: { //
    // non-critical section
    await !flag;
    flag = true;
    // CRITICAL SECTION
    flag = false;
    // non-critical section
}

thread Q: {
    // non-critical section
    await !flag;
    flag = true;
    // CRITICAL SECTION
    flag = false;
    // non-critical section
}
```

- ▶ Mutex: No
- ▶ Absence deadlock: Yes
- ▶ Free from starvation: Yes

Attempt 1

```
global int turn = 1;

1 thread P: {                                1 thread Q: {
2   // non-critical section                2   // non-critical section
3   await (turn==1);                      3   await (turn==2);
4   // CRITICAL SECTION                   4   // CRITICAL SECTION
5   turn = 2;                             5   turn = 1;
6   // non-critical section                6   // non-critical section
7 }                                         7 }
```

- ▶ turn is a “permission resource”
- ▶ Mutex:
- ▶ Absence deadlock:
- ▶ Free from starvation:

Attempt 1

```
global int turn = 1;

1 thread P: {                                1 thread Q: {
2   // non-critical section                  2   // non-critical section
3   await (turn==1);                        3   await (turn==2);
4   // CRITICAL SECTION                    4   // CRITICAL SECTION
5   turn = 2;                              5   turn = 1;
6   // non-critical section                6   // non-critical section
7 }                                         7 }
```

- ▶ turn is a “permission resource”
- ▶ Mutex: Yes
- ▶ Absence deadlock:
- ▶ Free from starvation:

Attempt 1

```
global int turn = 1;

1 thread P: {                                1 thread Q: {
2   // non-critical section                  2   // non-critical section
3   await (turn==1);                        3   await (turn==2);
4   // CRITICAL SECTION                     4   // CRITICAL SECTION
5   turn = 2;                               5   turn = 1;
6   // non-critical section                  6   // non-critical section
7 }                                          7 }
```

- ▶ turn is a “permission resource”
- ▶ Mutex: Yes
- ▶ Absence deadlock: Yes
- ▶ Free from starvation:

Attempt 1

```
global int turn = 1;

1 thread P: {                                1 thread Q: {
2   // non-critical section                  2   // non-critical section
3   await (turn==1);                        3   await (turn==2);
4   // CRITICAL SECTION                     4   // CRITICAL SECTION
5   turn = 2;                               5   turn = 1;
6   // non-critical section                  6   // non-critical section
7 }                                           7 }
```

- ▶ turn is a “permission resource”
- ▶ Mutex: Yes
- ▶ Absence deadlock: Yes
- ▶ Free from starvation: No

Attempt 1

```
global int turn = 1;

1 thread P: {                                1 thread Q: {
2   // non-critical section                  2   // non-critical section
3   await (turn==1);                        3   await (turn==2);
4   // CRITICAL SECTION                     4   // CRITICAL SECTION
5   turn = 2;                               5   turn = 1;
6   // non-critical section                  6   // non-critical section
7 }                                          7 }
```

- ▶ turn is a “permission resource”
- ▶ Mutex: Yes
- ▶ Absence deadlock: Yes
- ▶ Free from starvation: No (a process could remain indefinitely in its non-critical region)

Attempt 1

```
global int turn = 1;

1 thread P: {                                1 thread Q: {
2   // non-critical section                  2   // non-critical section
3   await (turn==1);                        3   await (turn==2);
4   // CRITICAL SECTION                    4   // CRITICAL SECTION
5   turn = 2;                              5   turn = 1;
6   // non-critical section                6   // non-critical section
7 }                                         7 }
```

- ▶ turn is a “permission resource”
- ▶ Mutex: Yes
- ▶ Absence deadlock: Yes
- ▶ Free from starvation: No (a process could remain indefinitely in its non-critical region)

How do we know mutex and absence of deadlock hold for sure?

Attempt I

- ▶ Mutex: Yes
- ▶ Absence deadlock: Yes

How do we know these properties hold for sure?

- ▶ We can resort to state diagrams or prove it using deductive systems
- ▶ For now we choose the former
- ▶ Let's build the state diagram

Attempt 1

```
global int turn = 1;

1 thread P: {
2     while (true) {
3         // non-critical section
4         await (turn==1);
5         // CRITICAL SECTION
6         turn = 2;
7         // non-critical section
8     }
9 }

1 thread Q: {
2     while (true) {
3         // non-critical section
4         await (turn==2);
5         // CRITICAL SECTION
6         turn = 1;
7         // non-critical section
8     }
9 }
```

Its meant to be executed in a loop

Attempt 1

```
global int turn = 1;

1 thread P: {                1 thread Q: {
2     while (true) {          2     while (true) {
3         await (turn==1);    3         await (turn==2);
4         turn = 2;           4         turn = 1;
5     }                      5     }
6 }                          6 }
```


Before drawing the state diagram, we get rid of irrelevant commands: we are only interested in synchronization

Attempt 1

```
global int turn = 1;


1 thread P: {
2   while (true) {
3     await (turn==1);
4     turn = 2;
5   }
6 }

1 thread Q: {
2   while (true) {
3     await (turn==2);
4     turn = 1;
5   }
6 }
```



Proving the Desired Properties I


- Mutex: Holds if all accesible states do no contain a state of the form $(p4, q4, turn)$ for some value of turn



Proving the Desired Properties II

Absence of deadlock (if some processes are trying to enter their CR, i.e. executing an await, then one must eventually succeed)

- Recall that the awaits are on line 3



Proving the Desired Properties III

Freedom from starvation (if any process is about to execute its preprotocol, then eventually it will succeed in entering the CR)

- Consider what happens if Q is trying to access its critical section but $turn = 1$ and P fails or loops in its NCS (*)

```
global int turn = 1;
```

```
1 thread P: {  
2   while (true) {  
3       // non-critical section (*)  
4       await (turn==1);  
5       // CRITICAL SECTION  
6       turn = 2;  
7       // non-critical section  
8   }  
9 }
```

```
1 thread Q: {  
2   while (true) {  
3       // non-critical section  
4       await (turn==2);  
5       // CRITICAL SECTION  
6       turn = 1;  
7       // non-critical section  
8   }  
9 }
```


Attempt 1 – Assessment

- ▶ Starvation is due to the fact that both processes test and set a single global variable
- ▶ If a process dies, the other is blocked
- ▶ Let us try to give each process its own variable

Attempt II

```
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        await !wantQ;
        wantP = true;
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        await !wantP;
        wantQ = true;
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

► Mutex:

Attempt II

```
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        await !wantQ;
        wantP = true;
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        await !wantP;
        wantQ = true;
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

► Mutex: No

Attempt III

```
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        await !wantQ;
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        await !wantP;
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

- ▶ Mutex:
- ▶ Absence deadlock:
- ▶ Free from starvation:

Attempt III

```
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        await !wantQ;
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        await !wantP;
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

- ▶ Mutex: Yes
- ▶ Absence deadlock:
- ▶ Free from starvation:

Attempt III

```
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        await !wantQ;
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        await !wantP;
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

- ▶ Mutex: Yes
- ▶ Absence deadlock: No
- ▶ Free from starvation:

Attempt III

```
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        await !wantQ;
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        await !wantP;
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

- ▶ Mutex: Yes
- ▶ Absence deadlock: No
- ▶ Free from starvation: No

Attempt III – Assessment

```
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        await !wantQ;
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        await !wantP;
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

- ▶ Process not only indicates its intention to enter but also insists on doing so
- ▶ Should:
 1. either back out if it discovers that it is contending with the other (Naive back-out);
 2. take turns (Dekker's algorithm);
 3. give priority to the first one that wanted access (Peterson's algorithm).
- ▶ We all these options in turn

Attempt IV – Naive Back-Out

```
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        while wantQ {
            wantP = false;
            wantP = true;
        }
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        while wantP {
            wantQ = false;
            wantQ = true;
        }
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

- ▶ Mutex:
- ▶ Absence deadlock:
- ▶ Free from starvation:

Attempt IV – Naive Back-Out

```
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        while wantQ {
            wantP = false;
            wantP = true;
        }
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        while wantP {
            wantQ = false;
            wantQ = true;
        }
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

- ▶ Mutex: Yes
- ▶ Absence deadlock:
- ▶ Free from starvation:

Attempt IV – Naive Back-Out

```
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        while wantQ {
            wantP = false;
            wantP = true;
        }
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        while wantP {
            wantQ = false;
            wantQ = true;
        }
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

- ▶ Mutex: Yes
- ▶ Absence deadlock: Yes
- ▶ Free from starvation:

Attempt IV – Naive Back-Out

```
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        while wantQ {
            wantP = false;
            wantP = true;
        }
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        while wantP {
            wantQ = false;
            wantQ = true;
        }
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

- ▶ Mutex: Yes
- ▶ Absence deadlock: Yes
- ▶ Free from starvation: No

Dekker's Algorithm (1965)

- ▶ Combines I and IV
- ▶ Attributed to Theodorus Jozef Dekker by Edsger W. Dijkstra in 1965

Dekker's Algorithm (I + IV)

```
global int turn = 1;
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        while wantQ
            if (turn == 2) {
                wantP = false;
                await (turn==1);
                wantP = true;
            }
        // CRITICAL SECTION
        turn = 2;
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        while wantP
            if (turn == 1) {
                wantQ = false;
                await (turn==2);
                wantQ = true;
            }
        // CRITICAL SECTION
        turn = 1;
        wantQ = false;
        // non-critical section
    }
}
```

Right to insist on entering is passed between the two processes

Peterson's Algorithm (1981)

```
global int last = 1;
global boolean wantP = false;
global boolean wantQ = false;

thread P: {
    while (true) {
        // non-critical section
        wantP = true;
        last = 1;
        await !wantQ or last==2;
        // CRITICAL SECTION
        wantP = false;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        wantQ = true;
        last = 2;
        await !wantP or last==1;
        // CRITICAL SECTION
        wantQ = false;
        // non-critical section
    }
}
```

Similar to Dekker except that if both want access, priority is given to the first one that wanted to access

Dekker and Peterson

- ▶ Mutex: Yes
- ▶ Absence deadlock: Yes
- ▶ Free from starvation: Yes

Race Conditions, Atomicity

Critical Sections

An Advanced Algorithm

The Bakery Algorithm

- ▶ Developed by Leslie Lamport
- ▶ Think of a Bakery shop
 - ▶ People take a ticket from a machine
 - ▶ If nobody is waiting, tickets don't matter
 - ▶ When several people are waiting, ticket order determines order in which they can make purchases
- ▶ Process interested in entering critical section acquires a ticket whose value is greater than all outstanding tickets
- ▶ Waits until its ticket is lowest value of all outstanding tickets

The Bakery Algorithm – Two Processes

```
global int np,nq =0;

thread P: {
    while (true) {
        // non-critical section
        np = nq + 1;
        await nq==0 or np<=nq;
        // CRITICAL SECTION
        np = 0;
        // non-critical section
    }
}

thread Q: {
    while (true) {
        // non-critical section
        nq = np + 1;
        await np==0 or nq<np;
        // CRITICAL SECTION
        nq = 0;
        // non-critical section
    }
}
```

- ▶ np and nq hold ticket numbers for each process
- ▶ Assumes assignment is atomic, which is not realistic
- ▶ We next take care of this and extend it to n processes

The Bakery Algorithm

```
global int[] number = new int[n]; // {0,0,...0}

thread {
    while (true) {
        // non-critical section
        number[id] = 1 + maximum(number);
        for all other processes j
            await (number[j]=0) or (number[id]<<number[j])
        // CRITICAL SECTION
        number[id] = 0;
        // non-critical section
    }
}
```

`number[i]<<number[j]` abbreviates

`(number[i]<number[j]) or (number[i]==number[j] and i<j)`

- ▶ Are ticket numbers be unbounded?
- ▶ Assumes computing the maximum of an array is atomic
- ▶ Must get rid of this otherwise mutual exclusion fails (why?)

The Bakery Algorithm

```
global boolean[] choosing = new boolean[n]; // {false,...}
global int[] number = new int[n]; // {0,0,...0}

thread {
    // non-critical section
    choosing[id] = true;
    number[id] = 1 + maximum(number);
    choosing[id] = false;
    for (j : range(0,n)) {
        await !choosing[j];
        await (number[j] == 0) or (number[id]<<number[j]);
    }

    // CRITICAL SECTION

    number[id] = 0;
    // non-critical section
}
```

The Bakery Algorithm

```
global boolean[] choosing = new boolean[n]; // {false,...}
global int[] number = new int[n]; // {0,0,...0}

thread {
    // non-critical section
    choosing[id] = true;
    number[id] = 1 + maximum(number);
    choosing[id] = false;
    for (j : range(0,n)) {
        await !choosing[j];
        await (number[j] == 0) or (number[id]<<number[j]);
    }

    // CRITICAL SECTION

    number[id] = 0;
    // non-critical section
}
```

This algorithm solves the problem for n threads.

Summary

- ▶ Difficult to solve the MEP using just atomic load and store
- ▶ If an atomic statement for both load and store were available, we could provide much easier solutions
- ▶ We'll see examples of this when we introduce such atomic statements next class