

Monitors

CS511

Review

- ▶ We've seen that **semaphores** are an efficient tool to solve synchronization problems
- ▶ However, they have some drawbacks
 1. They are low-level constructs
 - ▶ It is easy to forget an acquire or release
 2. They are not related to the data
 - ▶ They can appear in any part of the code

Monitors

- ▶ Combines ADTs and mutual exclusion
 - ▶ Proposed by Tony Hoare [1974]
- ▶ Adopted in many modern PLs
 - ▶ Java
 - ▶ C#

Main Ingredients

- ▶ A set of operations encapsulated in modules
- ▶ A unique **lock** that ensures mutual exclusion to all operations in the monitor
- ▶ Special variables called **condition variables**, that are used to program conditional synchronization

Counter Example

- ▶ Construct a counter with two operations: `inc()` and `dec()`.
- ▶ No two threads should be able to simultaneously modify the value of the counter
 - ▶ Think of a solution using semaphores
 - ▶ A solution using monitors

Counter using Semaphores

```
1 class Counter {
2     private int c = 0;
3     private Semaphore mutex = new Semaphore(1);
4
5     public void inc() {
6         mutex.acquire();
7         c++;
8         mutex.release();
9     }
10    public void dec() {
11        mutex.acquire();
12        c--;
13        mutex.release();
14    }
15 }
```

Counter using Monitors

```
1 monitor Counter {  
2  
3     private int counter = 0;  
4  
5     public void inc() {  
6         counter++;  
7     }  
8  
9     public void dec() {  
10        counter--;  
11    }  
12  
13 }
```

- ▶ The monitor comes equipped with a **lock or mutex** that allows at most one thread to execute its operations
- ▶ Note:
 - ▶ This is pseudocode (not Hydra nor Java)
 - ▶ We will see how to write monitors in Java later

Condition Variables

- ▶ Apart from the lock, there are **condition variables** associated to the monitor
- ▶ They have two operations:
 - ▶ `Cond.wait()`
 - ▶ `Cond.signal()`
- ▶ Just like in semaphores, they have an associated queue of blocked processes.

Condition Variables

`Cond.wait()`

- ▶ **always** blocks the process and places it in the waiting queue of the variable `Cond`.
- ▶ When it blocks, it releases the mutex on the monitor.

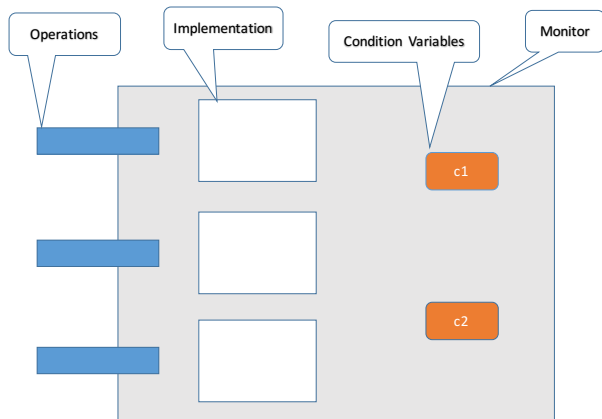
`Cond.notify()`

- ▶ Unblocks the first process in the waiting queue of the variable `Cond` and continues execution
- ▶ If there are no processes in the waiting queue, it has no effect.


Example: Buffer (Pseudocode)

```
1 monitor Buffer {
2     private condition spaceAvailable; // wait until space available
3     private condition dataAvailable;  // wait until data available
4
5     private Object data = null;      // shared data
6
7     public Object read() {
8         if (data == null)
9             dataAvailable.wait();
10        aux = data;
11        data = null;
12        spaceAvailable.signal();
13        return aux;
14    }
15
16    public void write(Object o) {
17        if (data != null)
18            spaceAvailable.wait();
19        data = o;
20        dataAvailable.signal();
21    }
22 }
23 }
```


Graphically




Typical Behavior



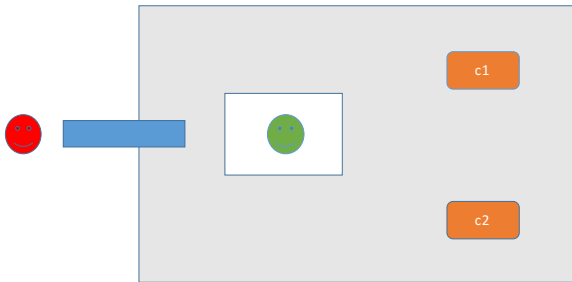
Typical Behavior




Typical Behavior




Typical Behavior



Typical Behavior




Wait




- ▶ Blocks process currently executing and associates it to variable's queue
- ▶ Upon blocking frees the **lock** allowing the entry of other processes

Wait




- ▶ Blocks process currently executing and associates it to variable's queue
- ▶ Upon blocking frees the **lock** allowing the entry of other processes

Wait



- ▶ Blocks process currently executing and associates it to variable's queue
- ▶ Upon blocking frees the **lock** allowing the entry of other processes

Signal




- ▶ Signalling process continues to execute after signalling on `c1`?
- ▶ Processes waiting in `c1`'s queue start immediately running inside the monitor?
- ▶ What about the processes blocked on entry to the monitor?


Signal (cont.)

States that a process can be in:

- ▶ Waiting to enter the monitor
- ▶ Executing within the monitor (only one)
- ▶ Blocked on condition variables
- ▶ Queue of processes just released from waiting on a condition variable
- ▶ Queue of processes that have just completed a `signal` operation



Signal



- ▶ Precedence of
 - ▶ signalling processes S
 - ▶ waiting processes W
 - ▶ processes blocked on entry E
- ▶ Two strategies:
 - ▶ **Signal and Urgent Wait:** $E < S < W$ (classical monitors, aka immediate resumption requirement – IRR)
 - ▶ **Signal and Continue:** $E = W < S$ (Java)

Signal and Urgent Wait

$$E < S < W$$

- ▶ When a process blocked on a condition variable is signaled, it immediately begins executing **ahead** of the signaling process
- ▶ Rationale: Signaling process changed the state of the monitor so that the condition now holds
- ▶ Waiting process resumes immediately (at the instruction immediately following the call to **wait** that blocked it) and hence need not check the condition
- ▶ Cons: signaling process unnecessarily delayed (unless signal is the last operation)

Examples

- ▶ Before analyzing the other strategy for signal, namely **Signal and Continue**, we present another example of a monitor
- ▶ We already considered the readers and writers problem with buffer of size 0
- ▶ We now consider the same problem with a buffer of size n
- ▶ Finally, we show how monitors can be used to define semaphores

Monitors

Examples

Monitors in Java


Exercise: N Dimensional Buffer

```
1 monitor Buffer {
2     private Object[] data = new Object[N+1];
3     private int begin = 0, end = 0;
4     private condition untilNotFull, untilAvailable;
5
6     public void push(Object o) {
7         if (isFull())
8             untilNotFull.wait();
9         data[begin] = o;
10        begin = next(begin);
11        untilAvailable.notify()
12    }
13
14    public Object pop() {
15        if (isEmpty())
16            untilAvailable.wait();
17        Object result = data[end];
18        end = next(end);
19        untilNotFull.notify();
20        return result;
21    }
22
23    private boolean isEmpty() { return begin == end; }
24    private boolean isFull()   { return next(begin) == end; }
25    private int next(int i)    { return (i+1)%(N+1); }
26
```

Monitor that Defines a Semaphore


```
1 monitor Semaphore {
2     private condition noCero;
3     private int permissions;
4
5     public Semaphore(int n) {
6         this.permissions = n;
7     }
8
9     public void acquire() {
10         if (permissions == 0)
11             noCero.wait();
12         permissions--;
13     }
14
15     public void release() {
16         permissions++;
17         noCero.signal();
18     }
19
20 }
```

Strategies for Signal



- ▶ Recall from above:
 - ▶ Precedence of signalling processes S
 - ▶ Precedence of waiting processes W
 - ▶ Precedence of processes blocked on entry E
- ▶ Two strategies for signal:
 - ▶ Signal and Urgent Wait: $E < S < W$ (classical monitors, aka immediate resumption requirement – IRR)
 - ▶ Signal and Continue: $E = W < S$ (Java)
- ▶ We already considered Signal and Urgent Wait
- ▶ We now consider Signal and Continue

Signal and Continue: $E = W < S$ (Java)



- ▶ Process from S which executes the signal continues execution
- ▶ Process from W which is unblocked joins competition for the lock
- ▶ Problem: signaling process can modify the condition after it executed the signal

Signal and Continue

- ▶ Must re-check the condition

```
public void acquire() {  
    while (permissions == 0)  
        noCero.wait();  
    permissions--;  
}
```

- ▶ Risk: introduce starvation.
- ▶ Seems less intuitive, but is preferred discipline today
 - ▶ Simpler formal semantics
 - ▶ Compatible with priority policies

Signal and Continue

```
1 public void acquire() {  
2     while (permissions == 0)  
3         noCero.wait();  
4     permissions--;  
5 }  
6  
7 public void release() {  
8     permissions++;  
9     noCero.signal();  
10 }
```

Signal and Continue

```
1 public void acquire() {  
2     while (permissions == 0)  
3         noCero.wait();  
4     permissions--;  
5 }  
6  
7 public void release() {  
8     permissions++;  
9     noCero.signal();  
10 }
```

- ▶ Is it fair?
- ▶ What happens with a process that is waiting to acquire the lock?

Signal and Continue

- ▶ Not fair because a process outside the monitor could steal the permission
- ▶ Possible measure: pass the permission on to a blocked process on the condition variable
- ▶ This requires that the process that executes the signal detects whether there are blocked processes on the associated condition variable

Monitor that defines a Semaphore

```
1 public void acquire() {
2     if (permissions == 0) {
3         waiting++;
4         noCero.wait();
5         waiting--;
6     } else {
7         permissions--;
8     }
9 }
10
11 public void release() {
12     if (waiting == 0)
13         permissions++;
14     else // else case does not increment permissions
15         noCero.signal();
16 }
```

Monitors

Examples

Monitors in Java

Monitors in Java

- ▶ Every class has a lock and a unique **condition variable**
 - ▶ Pros: convenient encapsulation (lock and condition variable cannot be tampered with)
 - ▶ Cons: Using multiple condition variables may benefit efficiency
- ▶ Methods `wait`, `notify` and `notifyAll` belong to the class `Object`
- ▶ Must use `synchronize` keyword in each method of the monitor
 - ▶ Guarantees mutual exclusion
 - ▶ Allows operations on condition variables to be invoked

N Dimensional Buffer in Java

```
1 class Buffer {
2     private Object[] data = new Object[N+1];
3     private int begin = 0, end = 0;
4
5     public synchronized void push(Object o) {
6         while (isFull()) wait();
7         data[begin] = o;
8         begin = next(begin);
9         notifyAll();
10    }
11
12    public synchronized Object pop() {
13        while (isEmpty()) wait();
14        Object result = data[end];
15        end = next(end);
16        notifyAll();
17        return result;
18    }
19
20    private boolean isEmpty() { return begin == end; }
21    private boolean isFull()  { return next(begin) == end; }
22    private int next(int i)   { return (i+1)%(N+1); }
23
24 }
```